

Florin Olariu and Lenuța Alboaie
Department of Computer Science
Alexandru Ioan Cuza University
Iași, Romania
florin.olariu@info.uaic.ro; lenuta.alboaie@info.uaic.ro

September 2026

Prof. Song Guo, Editor-in-Chief
IEEE Transactions on Cloud Computing
IEEE Computer Society

Dear Prof. Guo,

Please find enclosed our revised manuscript, TCC-2026-04-0292,

**Transparent Multi-Cloud Migration Planning:
A CSP–Expert–Pareto Decision Pipeline**

by Florin Olariu and Lenuța Alboaie, resubmitted in response to the Major Revision decision. We thank the two reviewers for their careful and constructive comments, which have materially improved the paper. A detailed, point-by-point response to every comment is provided in the accompanying “Response to Reviewers” document; below we summarize the main changes.

Summary of the revision.

- **No overstated optimizer claims.** Following Reviewer 1 (Comment 1) and Reviewer 2, we no longer imply that CMO is a stronger optimizer than the baselines on individual metrics. We report plainly that a tuned Genetic Algorithm can match or exceed CMO on raw cost and latency, and reposition CMO’s contribution around interpretability, auditability, and multi-criteria decision support rather than raw optimization performance.
- **Fair, like-for-like comparison.** Following Reviewer 1 (Comment 2) and Reviewer 2, all baselines are now also scored with the same six-dimensional MCDA scheme used to rank CMO’s solutions (cost, latency, reliability, security, vendor risk, scalability), so every comparison in the paper states explicitly which objective is being compared.
- **Honest scope on interpretability claims.** Following Reviewer 1 (Comment 3), we no longer imply that auditability and vendor-lock-in mitigation have been validated with users. We added a concrete, verbatim worked example of the constraint-proof artefact (Section V) and, drawing on Doshi-Velez and Kim’s established taxonomy of interpretability evaluation, we explain in the Response to Reviewers why the present functionally-grounded evidence is the appropriate bar for this revision and why a small, non-expert pilot would not properly answer the application-grounded question being asked; a properly resourced practitioner study is scoped as concrete future work. The manuscript itself makes no unsup-

ported claim in this regard.

- **Baseline configuration details.** Following Reviewer 1 (Comment 4), we removed an unexecuted, placeholder NSGA-II/MOEA-D comparison that had been included in error, and we now document the Genetic Algorithm’s encoding, population size, and operator settings in full so the remaining evolutionary baseline is reproducible.
- **Streamlined presentation.** Following Reviewer 1 (Comment 5), the two internal configurations formerly labelled CMOv3 and CMOv4 are now presented explicitly as two configurations of a single pipeline (differing only in how latency is modelled), not as sequential software releases, to keep the focus on the research contribution rather than on version history.
- **Condensed abstract.** Following Reviewer 1 (Comment 6), the abstract has been shortened by nearly half and no longer lists specific run counts, timings, or component counts, focusing instead on motivation, design, and the main qualitative findings.
- **Clarified feasibility vs. completeness.** Following Reviewer 2, we now state explicitly that the constraint layer guarantees the feasibility of every returned solution, and we clarify the conditions under which the reported Pareto set is, or is not, claimed to be complete.
- **Weight sensitivity.** Following Reviewer 2, we added a new analysis (Section IV-D) showing how the final MCDA ranking responds to changes in the stakeholder weight profile, together with the rationale for the default weights used elsewhere in the paper.
- **Modest framing of the Netflix case study.** Following Reviewer 2, the Netflix-inspired scenario is now presented as an illustrative, publicly reproducible example grounded in public documentation and pricing, not as validation against real production constraints.

Reproducibility. The implementation, experiment scripts, and result artefacts referenced in the manuscript are maintained in a repository that is now publicly available, as cited in the bibliography.

Declarations. We confirm that the revised manuscript is original, has not been published previously, and is not under consideration by any other journal or conference. All authors have approved the submitted version, and there are no conflicts of interest to declare.

We thank you and the reviewers again for the thorough evaluation, and we believe the revised manuscript addresses all of the points raised. Please let us know if any further

information or clarification is required.

Yours sincerely,

Florin Olariu (Corresponding Author)

Lenuța Alboaie (Corresponding Author)

Department of Computer Science, Alexandru Ioan Cuza University, Iași, Romania

`florin.olariu@info.uaic.ro; lenuta.alboaie@info.uaic.ro`